

# Introduction to Machine Learning and Deep Learning with Python

## Day 2

Tabea Rebafka



CNRS Formation Entreprise  
28 – 30 June, 2023, Paris

- 1 Decision trees and CART framework
- 2 Bagging
- 3 Random forests
- 4 Boosting – AdaBoost and Gradient boosting

- 1 Decision trees and CART framework
- 2 Bagging
- 3 Random forests
- 4 Boosting – AdaBoost and Gradient boosting

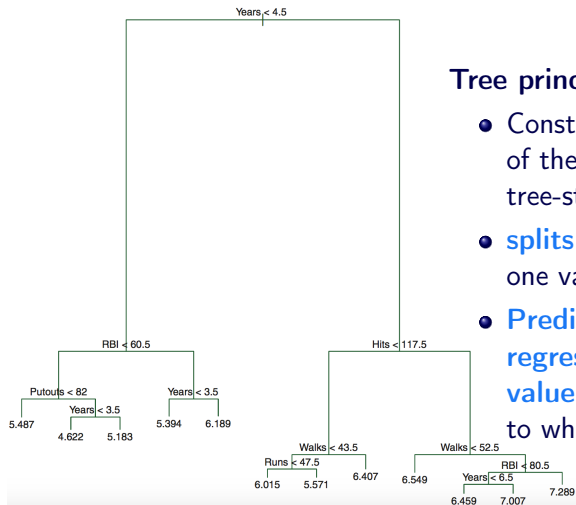
## Context

- For each individual  $i = 1, \dots, n$ , we observe
  - ▶ a **feature vector**  $x_i \in \mathbb{R}^d$  with  $x_i = (x_i^1, \dots, x_i^d) = (x_i^j)_{1 \leq j \leq d}$
  - ▶ an **outcome**  $y_i$  with values
    - ★ in  $\mathbb{R}$  (**regression problem**)
    - ★ in a finite number of classes  $\{1, \dots, K\}$  (**classification problem**)
- Data  $\mathcal{D} = \{(x_i, y_i), i = 1, \dots, n\}$  are considered to be i.i.d realizations of some random  $(X, Y)$
- We want to **learn the relation** between  $Y$  and  $X$
- Given a new feature vector  $x_{\text{new}}$ , we want to **predict** its unobserved outcome  $y_{\text{new}}$ , which is either a real value or a class label.

## Tree-based methods

- **stratification** or **segmentation** of the predictor space into a number of simple regions
- splitting rules used to segment the predictor space can be summarized in a tree called **decision tree**
- **prediction** are based on the values of the training data in the region to which the new observation  $x_{\text{new}}$  belongs

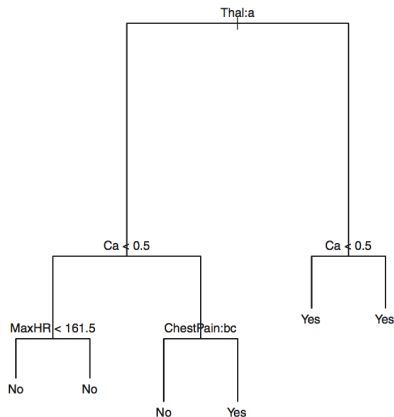
## Regression tree



## Tree principle

- Construct a **recursive partition** of the predictor space using a tree-structured set of “questions”
- **splits** around a given value of one variable  $x^j$  ( $1 \leq j \leq d$ )
- **Prediction** for  $x_{\text{new}}$  in the **regression** context: **mean value** of training data in the leaf to which  $x_{\text{new}}$  belongs

## Classification tree



## Tree principle

- Same construction
- **Prediction** for  $x_{\text{new}}$  in the **classification** context: **majority vote** in the leaf to which  $x_{\text{new}}$  belongs

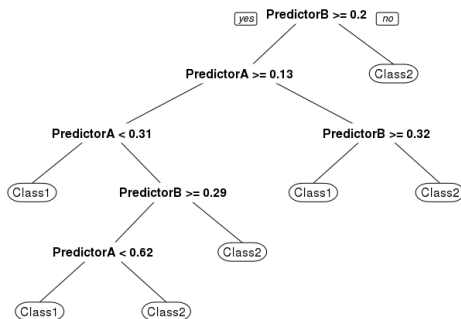

## Decision trees are

- a **simple** method
- highly **interpretable** (however, beware of unstable results)
- but **not competitive** with the best supervised learning approaches
- **improvement**: combine multiple trees to a single consensus prediction, called **ensemble methods** as
  - ▶ bagging
  - ▶ random forests
  - ▶ boosting



## CART = Classification and Regression Trees



- General framework for both types of decision trees



## Problem

- How to **construct a decision tree** with high prediction accuracy?
- Prediction quality depends on the tree (the partition) and there are many, many possible trees
- Finding the optimal tree is hard!

# Construction of decision trees II



values in the leafs/rectangles = proportion of **blue** training observations in the leaf/rectangle


## Construction of a tree

- **Top-down** construction where branches are created (**branching**)
- Start from a single region containing all training data (called **root** of the tree)
- Recursively choose a feature and a threshold to split nodes into ever smaller rectangular regions
- Leads to a **guillotine partition** of the feature space
- Which is equivalent to a **binary tree**

## How to split?

- Split such that the labels of the training data in the two new regions are as **homogeneous** as possible
  - ▶ For classification: the majority of observations in a node should have the same class label
  - ▶ For regression: the outcomes in a node should be very concentrated around their mean value

## Quantification of homogeneity




## Regression context

- **Variance** of training observations in node  $N \subset \{1, \dots, n\}$ :

$$V(N) = \frac{1}{\#N} \sum_{i \in N} (y_i - \bar{y}_N)^2$$

$$\text{with } \bar{y}_N = \frac{1}{\#N} \sum_{i \in N} y_i$$

## Quantification of homogeneity




## Regression context

- **Variance** of training observations in node  $N \subset \{1, \dots, n\}$ :

$$V(N) = \frac{1}{\#N} \sum_{i \in N} (y_i - \bar{y}_N)^2$$

$$\text{with } \bar{y}_N = \frac{1}{\#N} \sum_{i \in N} y_i$$

## Quantification of homogeneity




## Classification context

- different measures of **node (im)purity**
- all based on **class frequencies** of training observations in node  $N$ : For every class  $k$

$$p_{N,k} = \frac{\#\{i \in N : y_i = k\}}{\#N}$$



## Quantification of homogeneity



### Classification context

- different measures of **node (im)purity**
- all based on **class frequencies** of training observations in node  $N$ : For every class  $k$

$$p_{N,k} = \frac{\#\{i \in N : y_i = k\}}{\#N}$$

## Impurity measures for classification

- **Classification error rate**: fraction of the observations in node  $N$  that do not belong to the most common class:

$$1 - \max_k p_{N,k}$$

↪ not sufficiently sensitive for growing a tree

- **Gini index**

$$G(N) = \sum_{k=1}^K p_{N,k}(1 - p_{N,k})$$

- **Entropy** or cross-entropy

$$H(N) = - \sum_{k=1}^K p_{N,k} \log_2(p_{N,k})$$


↪  $G(N)$  and  $H(N)$  are quite similar numerically

↪ CART with Gini index is the most used technique

## Node splitting

- Split a node  $N$  into a left and a right child node  $N_L$  and  $N_R$
- Child nodes depend on the chosen **feature**  $j$  and the **threshold**  $t$  for the split

## Many possible splits



### Evaluate the quality of a split

- For a possible split of node  $N$  into children  $N_L$  and  $N_R$ , evaluate the **impurity** of the resulting child nodes by

$$\frac{|N_L|}{|N|} I(N_L) + \frac{|N_R|}{|N|} I(N_R)$$

where  $I$  denotes

- ▶ the variance  $V$  in the regression context
  - ▶ the Gini index  $G$  or the entropy  $H$  for classification
- weighted sum of impurity of both child nodes

## Best split

- Among **all** feature-threshold pairs  $(j, t)$  (i.e. for any  $j \in \{1, \dots, d\}$  and  $t \in \mathbb{R}$ ), the best split **minimizes** the associated **impurity** measure
- Equivalently, the best split **maximizes** the **information gain** defined as

$$\text{IG}(j, t) = I(N) - \frac{|N_L(j, t)|}{|N|} I(N_L(j, t)) - \frac{|N_R(j, t)|}{|N|} I(N_R(j, t))$$

where  $I(N)$  is the impurity of node  $N$

## Recursive partitioning

- CART method builds the partition iteratively
- starting from the root
- until some stopping criterion is satisfied

---

### Algorithm 1 CART

---

```
1: while stopping criterion is not met do
2:   for every node  $N$  of the current tree do
3:     Find the best feature-threshold pair  $(j, t)$  that maximizes  $IG(j, t)$ 
4:     Create the new child nodes
5:   end for
6: end while
```

---

## Greedy approach

- Greedy algorithms make the optimal choice at each step
- No regret strategy on the choice of the splits
- No guarantee to find the optimum

## Recall: Training and test sets

- The tree is learned on a **training set**
- Split data  $\mathcal{D} = \{(y_i, x_i), i = 1, \dots, n\}$  into a training set and a test set:  $\mathcal{D}_{\text{train}} \cup \mathcal{D}_{\text{test}} = \mathcal{D}$ 
  - ▶ Use  $\mathcal{D}_{\text{train}}$  to learn the tree
  - ▶ Use  $\mathcal{D}_{\text{test}}$  to evaluate error metrics and performance measures



## Stopping criterions


- Common criterions
  - ▶ Maximum depth of the tree is attained
  - ▶ Maximum number of leaves is attained
  - ▶ Number of data points per leaf is below a given threshold
  - ▶ Impurity per leaf is small enough
  - ▶ Test error increases
- Early stopping is always critical

## Alternative to stopping criterion


- Grow a maximal tree, then do **pruning**
- Not optimal either

## Regression trees versus linear regression

Linear Regression Fits Linear Data Better than Regression Tree



Regression Tree Fits Non-Linear Data Better than Linear Regression




## Classification tree versus logistic regression




# Illustration III

Decision Tree, f-measure = 0.889780



Logistic Regression, f-measure = 0.922420



## Importance of a specific feature

- Use the learned tree to evaluate the importance of each feature for the prediction task
- Different variable importance measures exist
- Naïvely, a feature  $j$  that appears in the splits is important
- An overall summary of the importance of feature  $j$  is the sum of the information gains obtained over all splits based on feature  $j$
- Beware that this measure does not account for correlations among features

- 1 Decision trees and CART framework
- 2 Bagging**
- 3 Random forests
- 4 Boosting – AdaBoost and Gradient boosting

## Inconvenients of decision trees

- CART suffers from
  - ▶ low prediction accuracy
  - ▶ lack of robustness, high variance: new training datasets give rise to rather different trees
- Decision trees are **weak predictors**

⇒ Instead of looking for one optimal tree, grow many, many trees and combine them to a more accurate predictor

⇒ Substantial improvements of prediction accuracy can be obtained by aggregating many decision trees (**ensemble methods**)

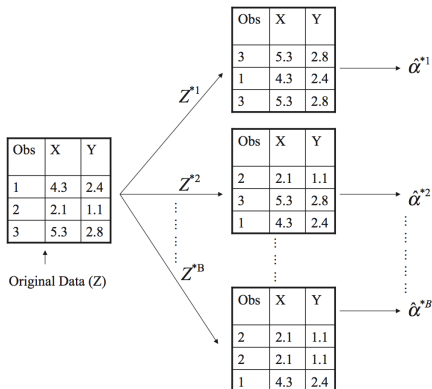
## Bagging approach

- General-purpose procedure to reduce the variance of **any statistical learning method**
- Bagging = Bootstrap Aggregation
- Particularly useful for decision trees
- Principle of **bagging**:
  - ▶ create many artificial datasets via **bootstrap**
  - ▶ learn many weak predictors on the bootstrapped samples
  - ▶ average/combine all predictors to obtain a highly accurate predictor
- **Variance reduction**: Let  $Z_1, \dots, Z_n$  be iid with variance  $\sigma^2$ . The variance of the mean  $\bar{Z}_n$  is  $\sigma^2/n$ . In other words, averaging a set of observations reduces variance.



## Bootstrap

- general **resampling** method to create new artificial datasets
- use bootstrapped samples to **evaluate the accuracy** (e.g. bias, variance, confidence interval) of any estimator or predictor



## Bootstrap datasets

- obtained by sampling **with replacement** from the original dataset: randomly choose observations  $(x_i, y_i)$
- generally contain some observations repeatedly
- same size as the original dataset

## General bootstrap approach

- Notations

- ▶  $\hat{f}$  : any statistical procedure or estimator that can be learned on a sample ;  $\hat{f}$  can be thought of an approximation of some “true” value or the “best” method  $f$
- ▶  $\hat{f}_{\mathcal{D}}$ : estimator  $\hat{f}$  trained on the original data  $\mathcal{D}$
- ▶  $\mathcal{D}^{(b)}$ ,  $b = 1, \dots, B$ : bootstrap datasets
- ▶  $\hat{f}^{(b)}$ : estimator learned on  $\mathcal{D}^{(b)}$  (**bootstrap replicates**)

- **The bootstrap replicates  $\{\hat{f}^{(b)}\}_{1 \leq b \leq B}$  are a sample of a distribution that is close to the distribution of  $\hat{f}$**
- Denote  $\bar{f}_B = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}$  the bootstrap mean
- Bootstrap approximation of the **bias**  $\mathbb{E}(\hat{f}) - f$  and the **variance** of  $\hat{f}$ :

$$\widehat{bias}_{\text{boot}}(\hat{f}) = \bar{f}_B - \hat{f}_{\mathcal{D}} \quad \widehat{var}_{\text{boot}}(\hat{f}) = \frac{1}{B} \sum_{b=1}^B \left( \hat{f}^{(b)} - \bar{f}_B \right)^2$$

---

## Algorithm 2 Bagging of regressors or classifiers

---

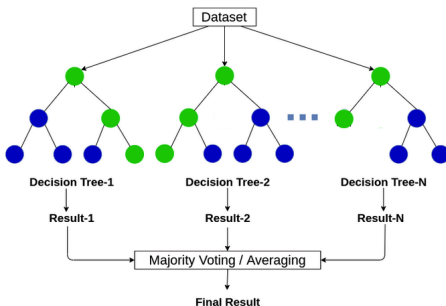
- 1: **for**  $b = 1, \dots, B$  **do**
  - 2:   Sample with replacement  $\mathcal{D}^{(b)}$  from the original dataset  $\mathcal{D}$
  - 3:   Fit a regressor (resp. classifier)  $\hat{f}^{(b)}$  on  $\mathcal{D}^{(b)}$
  - 4: **end for**
  - 5: Aggregate  $\hat{f}^{(1)}, \dots, \hat{f}^{(B)}$  to get a single regressor (resp. classifier) using an average (resp. majority vote)
- 

## Bagging for decision trees

- Trees are grown deep and not pruned. Each individual has high variance, but low bias.
- Aggregation of trees does not lead to a simple aggregated tree  
   $\hookrightarrow$  interpretability gets lost

## Prediction in a bagged model

- For a new point  $x_{\text{new}}$ , **every** bootstrap tree makes a prediction
- Aggregate all predictions to a single one:
  - ▶ in regression: average of all predictions
  - ▶ in classification: majority vote



## Out-of-bag error

- straightforward way to estimate the **test error** of a bagged model
- no cross validation, no validation set approach
- on average, a bootstrap sample only uses two-thirds of the data
- the remaining one-third are the **out-of-bag (OOB)** observations
- for every  $(x_i, y_i)$ , there are around one third of the bootstrap trees that were learned without using this observation

## Computation of the out-of-bag error

- For  $i = 1, \dots, n$ 
  - ▶ use all bootstrap trees that were learned when  $(x_i, y_i)$  was OOB and compute the prediction  $\hat{y}_i^{(b)}$  with the corresponding tree
  - ▶ average the predicted responses (or take a majority vote) to obtain a single prediction  $\hat{y}_i$
- use all predictions  $\hat{y}_i, i = 1, \dots, n$  to compute
  - ▶ in regression: the mean-squared error
  - ▶ in classification: a classification error

↪ The OOB error is a **valid estimate of the test error** for the bagged model

## Drawbacks of bagging

- Trees on bootstrap samples are relatively similar  $\hookrightarrow$  **correlated trees**
- If, for instance, the data contain a dominant feature for prediction, this feature appears at the top split of all trees
- Bagging has relatively **large variance**: when the whole bagging procedure is repeated on different training sets, quite different predictors are obtained

- 1 Decision trees and CART framework
- 2 Bagging
- 3 Random forests**
- 4 Boosting – AdaBoost and Gradient boosting



## Random forests

- Extension of bagging
- Aim: reduce variance
- Idea: force trees learned on bootstrap samples to be more different for a better exploration of the model space  $\hookrightarrow$  **decorrelate trees**

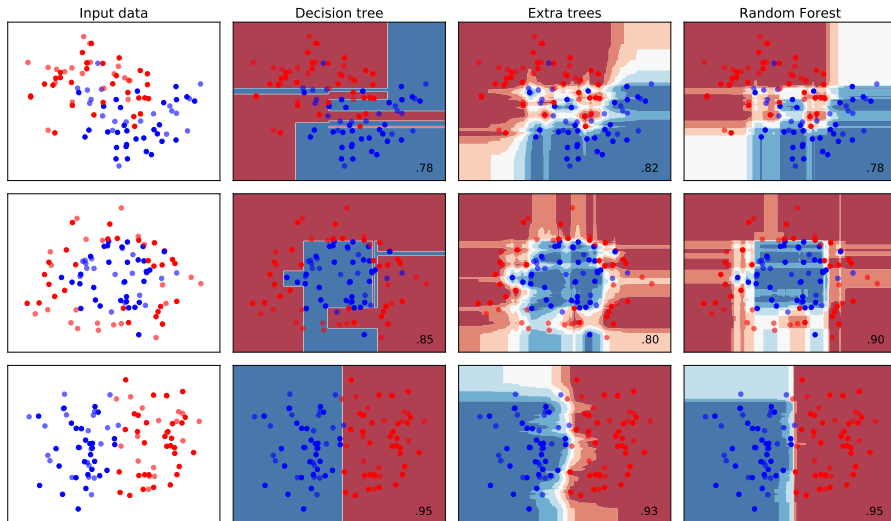
## Modification when growing a tree

- Do **column subsampling** also called **feature bagging**
- When looking for splits **only a random subset of features is tried**
- Reduces the correlation between trees, especially when some features are particularly strong

### Extremely randomized trees or Extra trees

- A variant of random forests and bagging
- Feature selection for the splits is done in the bagging fashion by minimizing an impurity measure (Gini, Entropy)
- **Thresholds are selected uniformly at random in the feature range**
- Computation is fast
- Leads to accurate predictors

# Random forests III



- 1 Decision trees and CART framework
- 2 Bagging
- 3 Random forests
- 4 Boosting – AdaBoost and Gradient boosting

## Ensemble learning


- Idea: combine many weak learners to form a strong one
- Bagging and random forests:
  - ▶ use independent learners
  - ▶ treat all weak learners equally
- Boosting:
  - ▶ weak learners are not independent
  - ▶ associate different weights to weak learners

## Boosting

- Sequential approach
- Each weak learner depends on the previously built learners
- Idea: correct for mistakes of the current learner
- Reweight the data points before feeding them to the model
- A simple reweighting approach: replicate certain observations in order to focus the classifier on predicting well on them


## AdaBoost

- = adaptive boosting
- for binary classification



- weak classifiers: vertical or horizontal half-planes
- $D_b$  : weights for data points ; start with equal weights

## First round



$$\epsilon_1 = 0.30$$


(weighted) error rate

$$\alpha_1 = 0.42$$


weight of the classifier in the ensemble



## Second round




## Third round



$$\epsilon_3 = 0.14$$

$$\alpha_3 = 0.92$$

## Final classifier



---

**Algorithm 3** AdaBoost

---

- 1: Initial weights of the data points:  $D_i^{(1)} = \frac{1}{n}$  for  $i = 1, \dots, n$
  - 2: **for**  $b = 1, \dots, B$  **do**
  - 3:   Train classifier:  $h_b \in \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n D_i^{(b)} 1_{y_i h(x_i) < 0}$
  - 4:   Error rate:  $\varepsilon_b = \sum_{i=1}^n D_i^{(b)} 1_{y_i h_b(x_i) < 0}$
  - 5:   Classifier weight:  $\alpha_b = \frac{1}{2} \log \left( \frac{1-\varepsilon_b}{\varepsilon_b} \right)$  (log odds ratio)
  - 6:   Normalizing constant:  $Z_b = 2\sqrt{\varepsilon_b(1-\varepsilon_b)}$
  - 7:   New data weights:  $D_i^{(b+1)} = D_i^{(b)} \frac{e^{-\alpha_b y_i h_b(x_i)}}{Z_b}$
  - 8: **end for**
  - 9: **return** Boosting classifier:  $g_B(x) = \text{sign} \left( \sum_{b=1}^B \alpha_b h_b(x) \right)$
-

## Boosting as a minimization algorithm

- Goal: an ensemble learner of the form

$$g_B(x) = \sum_{b=1}^B \alpha_b h_b(x)$$

that minimizes an empirical risk

$$\frac{1}{n} \sum_{i=1}^n \ell(y_i, g_B(x_i))$$

where  $\ell$  is a loss (least squares, logistic...)

- Minimization over a chosen family  $\mathcal{H}$  of weak learners and weights  $\alpha_b \in \mathbb{R}$

## Gradient boosting

- General-purpose method: many choices for the loss  $\ell$ , many choices of weak learners
- AdaBoost for binary classification:
  - ▶ exponential loss  $\ell(y, z) = e^{-yz}$
  - ▶ decision trees of depth 1, called stumps

## Problem

- How to find weak learners  $h_b$  and coefficients  $\alpha_b$ ?
- Even given  $\alpha_1, \dots, \alpha_B \in \mathbb{R}$ , minimize over  $|\mathcal{H}|^B$  to find the  $h_1, \dots, h_B$
- Size of  $\mathcal{H}$  is typically  $O(d)$  (number of features)

## Greedy algorithm

- At iteration  $b + 1$ : update current learner  $g_b$  by

$$g_{b+1} = g_b + \alpha_{b+1} h_{b+1}$$

where  $h_{b+1}$  and  $\alpha_{b+1}$  are solutions of

$$\arg \min_{\alpha \in \mathbb{R}, h \in \mathcal{H}} \sum_{i=1}^n \ell(y_i, g_b(x_i) + \alpha h(x_i))$$

## Still a problem

- Exact minimization at each step is too hard

## Gradient boosting idea

- Replace exact minimization by a gradient step
- Approximate the gradient step by an element of  $\mathcal{H}$

---

## Algorithm 4 Gradient boosting

---

- 1: Put  $g_0 = \arg \min_{m \in \mathbb{R}} \sum_{i=1}^n \ell(y_i, m)$
- 2: **for**  $b = 1, \dots, B$  **do**
- 3:   Get gradient directions:  $\delta_{b,i} = -\left(\nabla_u \ell(y_i, g_b(x_i) + u)\right)\Big|_{u=0}$
- 4:   Find learner  $h_{b+1}$  that approximates direction  $\delta_b$ :

$$(h_{b+1}, \nu_b) = \operatorname{argmin}_{h \in \mathcal{H}, \nu \in \mathbb{R}} \sum_{i=1}^n (\nu h(x_i) - \delta_{b,i})^2$$

- 5:   Classifier weight:  $\alpha_{b+1} = \arg \min_{\alpha \in \mathbb{R}} \sum_{i=1}^n \ell(y_i, g_b(x_i) + \alpha h_{b+1}(x_i))$
  - 6:   Update ensemble learner:  $g_b = g_{b-1} + \alpha_b h_b$
  - 7: **end for**
  - 8: **return** Boosting learner  $g_B$
-



## Hyperparameters

- Family  $\mathcal{H}$  of weak learners:
  - ▶ If  $\mathcal{H}$  contains decision trees, choose their depth (usually small... no more than four or five)
  - ▶ If  $\mathcal{H}$  contains generalized linear models (logistic regression), select the number of features (usually one or two)
- Number of iterations  $B$ :
  - ▶ overfitting is possible (contrary to bagging and random forests)
  - ▶ use cross validation, stop when test error does no longer improve
  - ▶ usually a few hundreds, few thousands
- Check documentation of `GradientBoostingRegressor` of the `sklearn.ensemble` module for more details on the hyperparameters

## Many implementations available

... with many clever extra tricks

- GradientBoostingRegressor

- ▶ Part of the `sklearn.ensemble` module

- XGBoost

- ▶ <https://xgboost.readthedocs.io/en/latest/>
- ▶ <https://github.com/dmlc/xgboost>
- ▶ Has been state of the art for years

- LightGBM

- ▶ <https://lightgbm.readthedocs.io/en/latest/>
- ▶ <https://github.com/Microsoft/LightGBM>
- ▶ Faster (and better ?) than XGBoost

- CatBoost

- ▶ <https://catboost.ai/en/docs/>
- ▶ One of the most recent implementations
- ▶ Adapted for categorical features

# Grand mother's recipe I



For standard tabular data (no image, signal, text etc.)

- Spend time on **feature engineering**
- Spend time on **feature engineering**
- ...
- Spend time on **feature engineering**
- Always try out random forests or gradient boosting **before** diving into a deep learning method (tomorrow)

## This afternoon: practical session

- Prediction of electricity demand based on load curves
- Using random forests and LightGBM

# Thank you!